

Sushant Mudalgi

Boston, Massachusetts | mudalgi.s@northeastern.edu | +1 (617) 407-3780 | linkedin.com/in/sushant-mudalgi | github.com/sushantm47

Data Engineer with experience building ETL/ELT pipelines and data models on Databricks and AWS, processing large-scale datasets for enterprise analytics and regulatory reporting.

Education

Northeastern University, Boston - MS in Computer Science Aug 2025 - May 2027

Coursework: Algorithms, Advanced Database Management Systems, Foundations of AI, Program Design Paradigms

KLE Technological University, Hubli - BE in Computer Science GPA: 3.91 / 4.00 Aug 2017 - May 2021

Coursework: Distributed Cloud Computing, Computer Networks, Database Management Systems, Machine Learning

Experience

Data Engineer L4, Takeda - Bengaluru, India May 2025 - Aug 2025

Joined directly by Takeda following account work on their data engineering team while at EY GDS.

- Developed and maintained 15+ production PySpark ETL pipelines on Databricks spanning North America and Europe regions, powering analytics and regulatory reporting for a global pharmaceutical enterprise under strict data-compliance requirements.
- Engineered secure data ingestion using Databricks Unity Catalog and AWS (S3, DMS, Step Functions, Lambda), including scheduled snapshot delivery from Databricks Delta Marketplace to guarantee timely, reliable data availability for downstream analytics.

Data Engineer, EY GDS - Bengaluru, India Nov 2024 - Apr 2025

- Implemented 15 real-time and batch PySpark pipelines processing 200-300 GB of data from heterogeneous sources (SQL Server, SFTP, SharePoint, Heroku, third-party vendors), writing modular, object-oriented Python and PySpark code for ingestion, validation, and transformation at scale.
- Designed 30+ Delta Lake tables optimized for high-performance queries and pharmaceutical reporting; automated ingestion and processing with AWS DMS, EventBridge, Step Functions, and S3.
- Migrated orchestration from AWS Step Functions to Tidal, cutting workflow costs by 25%, and scheduled 10+ Databricks ETL workflows with clearer job dependencies and improved reliability.

Software Engineering Analyst, Deloitte USI - Bengaluru, India Sep 2021 - Nov 2024

- Architected fact and dimension tables in Amazon Redshift and integrated the warehouse with Tableau, delivering dimensional models and dashboards that supported enterprise reporting and analytics.
- Built ETL jobs in Talend and AWS Glue, applying OOP principles to write reusable, maintainable Python transformations for cleaning transactional data, and led the migration of Talend jobs to AWS Glue, improving data-processing efficiency.
- Applied machine learning algorithms (K-Means, Mean Shift, Linear Regression, XGBoost) in Python to surface business insights from cleaned transactional datasets.

Software Engineering Intern (Web Development), Kalpas Innovations - Hubli, India Mar 2021 - Jul 2021

- Created the front end of a digital platform in React.js with real-time user interactions, optimized for 5,000+ monthly active users, and built a document-oriented MongoDB store for efficient data storage and retrieval.

Projects

DealRadar: US Price Monitoring Dashboard | Python, Streamlit, Pandas, PostgreSQL, BeautifulSoup 2025

- Delivered a live price-monitoring dashboard in Streamlit tracking product prices and availability across Amazon US and Walmart US, developing Python web-scraping pipelines (Requests, BeautifulSoup4) with structured extraction of pricing data.
- Modeled and stored price history in PostgreSQL for real-time querying and analytics; designed the dashboard to surface lowest prices, architecting for planned extensions including price-drop alerting and a RAG-based product-insights assistant over price history.

Technical Skills

Languages and Frameworks: Python, SQL, JavaScript, PySpark, Spark (DataFrames, Spark SQL, Datasets), FastAPI, React.js, Tableau

Software Engineering: Data Structures and Algorithms, Object-Oriented Design, REST API Development (FastAPI), Git/GitHub, CI/CD (Jenkins, OpenShift), Agile/Scrum

Data Engineering and Distributed Systems: ETL/ELT Pipeline Design, Real-Time and Batch Processing, Dimensional Modeling (Star Schema), Data Validation and Quality, Workflow Orchestration (Tidal, Databricks Workflows)

Cloud and Platforms: AWS (S3, Glue, Redshift, Lambda, DMS, EventBridge, Step Functions), Databricks (Delta Lake, Unity Catalog, Delta Marketplace, Azure Databricks)

Databases, ML, and Tools: SQL Server, PostgreSQL, MongoDB, Talend, Splunk, scikit-learn, XGBoost, K-Means, Linear Regression, Pandas, NumPy